

# CS458: Introduction to Information Security

## Notes 9: Key Establishment

Yousef M. Elmehdwi

Department of Computer Science

Illinois Institute of Technology

[yelmehdwi@iit.edu](mailto:yelmehdwi@iit.edu)


October 25, 2018

Slides: Modified from [Christof Paar and Jan Pelzl](#) & [Ewa Syta](#)

- Introduction
- The  $n^2$  Key Distribution Problem
- Symmetric Key Distribution
- Asymmetric Key Distribution
  - Man-in-the-Middle Attack
  - Certificates
  - Public-Key Infrastructure

# Classification of Key Establishment Methods

- **key establishment** deals with establishing a shared secret between two or more parties
- Methods for this can be classified into **key transport** and **key agreement** methods




- In an ideal key agreement protocol, no single party can control what the key value will be.
- Additional variations beyond key transport and key agreement exist, including various forms of key update, such as key derivation

# Key Freshness

- It is often desirable to frequently change the key in a cryptographic system.
- Reasons for key freshness include:
  - If a key is exposed (e.g., through hackers), there is limited damage if the key is changed often
  - Some cryptographic attacks become more difficult if only a limited amount of ciphertext was generated under one key
  - If an attacker wants to recover long pieces of ciphertext, he has to recover several keys which makes attacks harder.

# Key Derivation

- In order to achieve **key freshness**, we need to generate new keys frequently.
- Rather than performing a full key establishment every time (which is costly in terms of computation and/or communication), we can derive multiple **session keys**  $k_{ses}$  from a given key  $k_{AB}$ .



- The key  $k_{AB}$  is fed into a **key derivation function (KDF)** together with a **nonce**  $r$  ("number used only once").
- Every different value for  $r$  yields a different **session key**

# Key Derivation

- The key derivation function is a computationally simple function, e.g., a block cipher or a hash function.
- it should be a one-way function
- Example for a basic protocol:


**Alice**

derive key

$$k_{ses} = e_{k_{AB}}(r)$$

**Bob**

generate nonce  $r$




derive key

$$k_{ses} = e_{k_{AB}}(r)$$

- Other alternatives:
  - hashing *nonce* together with  $k_{AB}$ :  $K_{ses} = \text{HMAC}_{k_{AB}}(r)$
  - Rather sending a *nonce*, Alice and Bob can also simply encrypt a counter *cnt* periodically:  $K_{ses} = e_{k_{AB}}(cnt)$  or
  - Compute the HMAC of the counter:  $K_{ses} = \text{HMAC}_{k_{AB}}(cnt)$

# Naïve approach: The $n^2$ Key Distribution Problem

- We assumed the necessary keys for symmetric algorithms are distributed via a “secure channel”
- **key predistribution** since it involves different mode of communication, e.g., the key is transmitted via a phone line, or in a letter.
- **Simple situation:** Network with  $n$  users. Every user wants to communicate securely with every of the other  $n-1$  users.
- **Naïve approach:** Every pair of users obtains an individual key pair
- Example: Keys in a network with  $n = 4$  users:



# The $n^2$ Key Distribution Problem

- Shortcomings

- There are  $n(n-1) \approx n^2$  keys stored in the system
- There are  $\frac{n(n-1)}{2}$  pair keys
- If a new user Sammy joins the network, new keys  $k_{XS}$  have to be transported via secure channels to each of the existing users  
 $\Rightarrow$  Only works for small networks which are relatively static
- Example: mid-size company with 750 employees pause  
 $750 \times 749 = 561,750$  keys must be distributed securely.



# Key Establishment with Key Distribution Center (KDC)

- **Key Distribution Center (KDC)**: Central party, trusted by all users
- **Idea**:
  - Central “trusted authority” (KDC) that shares one key **key encryption key (KEK)** with every user
  - KDC sends session keys to users which are encrypted with *KEK*s

**Alice**

KEK:  $k_A$

**KDC**

KEK:  $k_A, k_B$

**Bob**

KEK:  $k_B$

$\xrightarrow{\text{RQST}(ID_A, ID_B)}$

generate random  $k_{ses}$

$y_A = e_{k_A}(k_{ses})$

$y_B = e_{k_B}(k_{ses})$

$\xleftarrow{y_A}$

$\xrightarrow{y_B}$

$k_{ses} = e_{k_A}^{-1}(y_A)$

$k_{ses} = e_{k_B}^{-1}(y_B)$

---

$y = e_{k_{ses}}(x)$

$\xrightarrow{y}$

$x = e_{k_{ses}}^{-1}(y)$

- Example: Use AES with *KEK*s and fast stream cipher with **short-lived**  $k_{ses}$

# Key Establishment with Key Distribution Center (KDC)

- Advantages over previous approach:
  - Only  $n$  **long-term key pairs** are in the system
  - If a new user is added, a secure key is only needed between the user and the KDC (the other users are not affected)
  - Scales well to moderately sized networks


# Key Establishment with Key Distribution Center (KDC)

- How does Bob know he is talking to Alice?
  - **Replay attack**: Eve breaks an old session key, resends old messages  $y_A$  and  $y_B$ , and can then decrypt all messages
  - **Key confirmation attack**: Eve intercepts Alice's request  $RQST(ID_A, ID_B)$  and sends to KDC  $RQST(ID_A, ID_e)$ . Alice believes that she communicates with Bob and Eve will decrypt her messages - there is **no key confirmation**
- Kerberos (a popular authentication and key distribution protocol) is based on KDCs

# Kerberos - Authentication and Key Distribution Standard Protocol

- Advanced protocol that protects against both replay and key confirmation attacks
- Its main purpose is to provide user authentication in computer networks
- Based on a KDC, which is named the “authentication sever” in Kerberos terminology.

# Key Establishment Using a Simplified Version of Kerberos




- Protects against both attacks
- Adds **Lifetime**  $T$ , **Time stamp**  $T_S$  and **Nonce** (only KDC can send new keys)

# Key Establishment with Key Distribution Center

- Remaining problems:
  - **No Perfect Forward Secrecy**: If the *KEK*s are compromised, an attacker can decrypt past messages if he stored the corresponding ciphertext
  - **Single point of failure**: The KDC stores all *KEK*s. If an attacker gets access to this database, all past traffic can be decrypted.
  - **Communication bottleneck**: The KDC is involved in every communication in the entire network (can be countered by giving the session keys a long life time)
  - **Secure channel during initialization**: when a new user joins - public key cipher for new key transport
- A cryptographic protocol has **perfect forward secrecy (PFS)** if the compromise of **long-term keys** does not allow an attacker to obtain **past session keys**
- The main mechanism to assure PFS is to employ public-key techniques.

## Key Establishment Using Asymmetric Techniques


# Recall: Diffie-Hellman Key Exchange (DHKE)



- Widely used in practice
- If the parameters are chosen carefully (especially a prime  $p > 2^{1024}$ ), the DHKE is secure against passive (i.e., listen-only) attacks
- However: If the attacker can actively intervene in the communication, the **man-in-the-middle attack** becomes possible.



# Man-in-the-Middle Attack



- Oscar computes a session key  $k_{AO}$  with Alice, and  $k_{BO}$  with Bob
- However, Alice and Bob think they are communicating with each other!
- The attack efficiently performs 2 DH key-exchanges: Oscar-Alice and Oscar-Bob
- Here is why the attack works:

Alice computes:  $k_{AO} = (B')^a = (\alpha^o)^a$       Bob computes:  $k_{BO} = (A')^b = (\alpha^o)^b$

Oscar computes:  $k_{AO} = A^o = (\alpha^a)^o$       Oscar computes:  $k_{BO} = B^o = (\alpha^b)^o$

- Oscar has now complete control over the channel, e.g., if Alice wants to send an encrypted message  $x$  to Bob, Oscar can read the message

# Important facts about the Man-in-the-Middle (MITM) Attack

- The man-in-the-middle-attack is not restricted to DHKE; it is applicable to any public-key scheme, e.g. RSA encryption, etc.
- Attack works always by the same pattern: Oscar replaces the public key from one of the parties by his own key.
- The attack is also known as Janus attack<sup>1</sup>
- **Q:** What makes the MITM attack possible?
  - The public keys are not authenticated: When Alice receives a public key which is allegedly from Bob, she has no way of knowing whether it is in fact his.
- *Even though public keys can be sent over insecure channels, they require authenticated channels*

---


<sup>1</sup>in reference to the roman two-headed god of gates

# Certificates


- In order to authenticate public keys (and thus, prevent the MIM attack), all public keys are digitally signed by a **central trusted authority**.
- Such a construction is called **certificate**  
**certificate = public key + ID(user) + digital signature over public key and ID**
- In its most basic form, a certificate for the key  $k_{pub}$  of user Alice is:  
$$Cert(Alice) = (k_{pub}, ID(Alice), sig_{K_{CA}}(k_{pub}, ID(Alice)))$$
- Certificates bind the identity of user to her public key
- The trusted authority that issues the certificate is referred to as **certifying authority (CA)**
- “Issuing certificates” means in particular that the CA computes the signature  $sig_{K_{CA}}(k_{pub})$  using its (super secret!) private key  $k_{CA}$ .
- The party who receives a certificate, e.g., Bob, verifies Alice’s public key using the public key of the CA

# Certificate generation

- User provided keys:



- CA generated keys:



# Diffie-Hellman Key Exchange (DHKE) with Certificates

**Alice**

$$a = k_{pr,A}$$

$$A = k_{pub,A} \equiv \alpha^a \bmod p$$

$$\text{Cert}_A = [(A, ID_A), s_A]$$

verify certificate:

$$\text{ver}_{k_{pub,CA}}(\text{Cert}_B)$$

compute session key:

$$k_{AB} \equiv B^a \bmod p$$

**Bob**

$$b = k_{pr,B}$$

$$B = k_{pub,B} \equiv \alpha^b \bmod p$$


$$\text{Cert}_B = [(B, ID_B), s_B]$$

verify certificate:

$$\text{ver}_{k_{pub,CA}}(\text{Cert}_A)$$

compute session key:

$$k_{AB} \equiv A^b \bmod p$$



# Certificates

- Note that verification requires the public key of the CA for  $ver_{k_{pub}, CA}$
- In principle, an attacker could run a MIM attack when  $k_{pub}, CA$  is being distributed  
 $\Rightarrow$  The public CA keys must also be distributed via an authenticated channel!
- Q: So, have we gained anything?  
After all, we try to protect a public key (e.g., a DH key) by using yet another public-key scheme (digital signature for the certificate)?
- YES! The difference from before (e.g., DHKE without certificates) is that we only need to distribute the public CA key once, often at the set-up time of the system
- Example: Most web browsers are shipped with the public keys of many CAs. The “authenticated channel” is formed by the (hopefully) correct distribution of the original browser software.

- Public Key Infrastructure (PKI): all you need to securely use public key crypto.

i.e., The entire system that is formed by CAs together with the necessary support mechanisms

- Key generation and management.
- Certificate authority (CA) or authorities or other trust model.
- Certificate revocation lists (CRLs), etc.

# Certificates in the Real World: X.509 Certificates

- Certificates contain more information than just a public key and a signature
- X.509 is an important standard for network authentication services, and the corresponding certificates are widely used for Internet communication, i.e., in S/MIME, IPsec and SSL/TLS

|   |
|---|
| Serial Number   |
| Certificate Algorithm: <ul style="list-style-type: none"><li>- Algorithm</li><li>- Parameters</li></ul>                     |
| Issuer  |
| Period of Validity: <ul style="list-style-type: none"><li>- Not Before Date</li><li>- Not After Date</li></ul>              |
| Subject   |
| Subject's Public Key: <ul style="list-style-type: none"><li>- Algorithm</li><li>- Parameters</li><li>- Public Key</li></ul> |
| Signature   |



# Certificates in the Real World: X.509 Certificates

- **Certificate Algorithm:** Specified which signature algorithm is being used, e.g., RSA with SHA-1 or ECDSA with SHA-2, and with which parameters, e.g., the bit lengths.
- **Issuer:** There are many companies and organizations that issue certificates. This field specifies who generated the one at hand.
- **Period of Validity:** public key is not certified indefinitely but rather for a limited time
- **Subject:** This field contains what was called  $ID_A$  or  $ID_B$  in our earlier examples. It contains identifying information such as names of people or organizations.
- **Subject's Public Key:** The public key that is to be protected by the certificate. In addition, the algorithm and its parameters are stored.
- **Signature:** The signature over all other fields of the certificate
- Note that there are two public-key schemes involved in every certificate: the one whose public key is protected by the certificate and the algorithm with which the certificate is signed

# Remaining Issues with PKIs

- There are additional problems when certificates are to be used in systems with a large number of participants. The more pressing ones are:
  1. Users communicate with others whose certificates are issued by different CAs
    - This requires cross-certification of CAs, e.g., CA1 certifies the public-key of CA2. If Alice trusts “her” CA1, cross-certification ensures that she also trusts CA2. This is called a “chain of trust” and it is said that “trust is delegated”.
  2. Certificate Revocation Lists (CRLs)
    - Another real-world problem is that certificates must be revoked, e.g., if a smart card with certificate is lost or if a user leaves an organization. For this, CRLs must be sent out periodically (e.g., daily) which is a burden on the bandwidth of the system.

- Key agreement vs. Key transport
- Key freshness
- Key derivation
- Key establishment can be done using symmetric or asymmetric (public key) techniques
- Key establishment with a Key Distribution Center
- Certificates